

MARKETMIND AI



User Guide

Running and using the MarketMind AI workspace

Contents

1. What this guide covers	3
2. Before you start (local setup)	3
3. Onboarding	3
4. The journey, step by step	4
5. Frequently asked questions	5
6. Troubleshooting	6

1 What this guide covers

How to run MarketMind AI locally, sign in, and take a business through the full journey: Discovery → Strategy → Content → owner review → export/publish.

2 Before you start (local setup)

Prerequisites: Node.js 20+, Docker with Docker Compose, and `uv` (Python 3.12+).

```
# 1. Install dependencies
npm install

# 2. Configure environment (one per app)
cp apps/api/.env.example apps/api/.env
cp apps/web/.env.example apps/web/.env
cp services/ai/.env.example services/ai/.env

# 3. Run the full stack (Postgres + Redis + API + AI + web), with migrations:
npm run dev:full
```

Then open `http://localhost:3000` in your browser.

Service	URL
Web (Next.js)	<code>http://localhost:3000</code>
API (NestJS)	<code>http://localhost:3001/api/v1</code>
AI (FastAPI)	<code>http://localhost:8000</code>
Health check	<code>GET http://localhost:3001/api/v1/health</code>

The `dev` command starts Docker (Postgres on `5433` , Redis on `6379`), applies Prisma migrations, and launches the three apps in parallel. Stop everything with `Ctrl+C` .

3 Onboarding

1. Open the app. Unprefixed URLs are redirected to your locale (`/ar/` by default, or your preferred language).
2. Switch language at any time. Arabic (RTL) and English are both supported; the interface flips direction correctly.
3. Sign in to create your owner account (Google OAuth is configured for local development).
4. The workspace home is ready to start Discovery.

For the graduation MVP there is one owner account and one business profile. Multi-business and agency dashboards are deferred.

4 The journey, step by step

Step 1 Discovery

The Discovery Agent runs a short bilingual interview, asking one clear question at a time about:

- business name and location,
- the product or service,
- target customers and goals,
- budget and brand style,
- social media accounts and competitors,
- optional uploads: menu, logo, brand guide, photos.

Review and confirm the resulting Business Profile before anything else can proceed. The profile records unknowns rather than inventing facts.

After you approve the profile, Strategy can begin.

Step 2 Research

The system collects trusted, cited marketing context: curated internal marketing knowledge first, limited web research only when needed. Each fact carries a source reference and a short note on why it matters. Sources remain visible to you.

Step 3 Strategy

The Strategy Agent produces a plan that you can approve, edit, or reject:

- primary goal and target audience,
- recommended platforms and tone/language,
- content themes,
- a four-week detailed plan and a twelve-week overview,
- budget direction and KPIs,
- citations, plus visible assumptions, knowledge gaps, and blockers.

After you approve the strategy, weekly content generation can start.

Step 4 Content

From the approved strategy, the system generates content one week at a time (3–5 posts per week):

- Arabic and English captions where useful,

- image ideas or generated assets,
- short-video scripts,
- posting notes.

Week 1 is created when the cycle starts; by the end of each week the next week's draft is ready for review. Automatic next-week generation creates drafts only; it never approves or publishes.

You approve individual items (or a whole week) before anything moves forward.

Step 5 Publish or export

Approved content can leave the system in one of three ways:

1. **Real publish** — through an approved, deterministic n8n/Meta workflow, only when the target account, permissions, and media are ready and you separately approve the exact candidate, target, mode, and schedule.
2. **Export** — a checksum-addressed content package you can use anywhere.
3. **Simulation** — a clearly labeled deterministic demo simulation, never presented as real.

APPROVAL BOUNDARY

Publishing is not an AI agent. It is a logged, deterministic action that always requires your explicit approval.

Step 6 Monitor (Facebook, read-only)

For posts MarketMind actually published, the system collects read-only Facebook Page Insights automatically at comparable ages (24h, 72h, 7 days). Test fixtures and simulations never appear as live analytics.

Step 7 Improve

After enough comparable data, the system may suggest a hook-style or CTA-wording change for one future draft. It cannot modify the approved strategy, weekly plan, topics, audience, channel, format, locale, or past content. You approve or dismiss it.

5 Frequently asked questions

Does the AI publish by itself?

No. Publishing is deterministic and gated on your explicit approval of the exact candidate, target, mode, and schedule.

Is my business data used to train anything?

No. Business profile data stays in your local PostgreSQL. Only reviewed, curated marketing knowledge is embedded into the shared Qdrant index, never your profile.

Why do I see "simulation"?

Where a real external integration is not yet connected for the demo, the system shows a clearly labeled simulation so you never mistake it for a real publish or real analytics.

Can I use Arabic?

Yes. Arabic is a first-class experience with correct RTL, and English is fully supported.

What platforms are supported for publishing?

Facebook (static posts) via an approved workflow. Instagram/TikTok publishing and paid ads are deferred for this submission.

Can I manage several businesses?

Not in this MVP. One owner, one business profile.

6 Troubleshooting

Symptom	Likely cause	Fix
<code>GET /api/v1/health</code> times out	Postgres/Redis not up	<code>npm run docker:up</code> , then re-run <code>npm run dev</code>
Prisma migration error	DB schema drift	<code>npm run dev:full</code> (applies <code>prisma migrate deploy</code>)
AI <code>/health</code> shows <code>qdrant: unreachable</code>	Qdrant not running	<code>npm run docker:strategy:up</code> (starts Postgres + Redis + Qdrant)
Strategy retrieval empty	Knowledge not ingested	<code>npm run strategy:knowledge:sync:local</code>
Image generation fails	<code>IMAGE_MODEL</code> missing <code>google/</code> prefix for OpenRouter	set <code>IMAGE_MODEL=google/gemini-3.1-flash-image</code> in <code>services/ai/.env</code>
Login loop / 401	JWT secrets rotated	sign in again (old tokens are invalid after a secret change)
Very slow first compile	repo on an external/network drive	move the repo to a local (ext4) filesystem, or wait for the initial compile

Screenshots are captured during the demo rehearsal (see the MarketMind AI Demo Runbook) into `assets/` , covering the landing and in-app steps in English and Arabic (RTL).